

# LEAD SCORING PROJECT

**\*\* lead scoring from business understanding to machine learning with pandas and scikit-learn or other required libraries!**

## Problem context

Lead scoring is a process of assigning scores to prospects based on their profile and behavioral data in order to prioritize leads, improve close rates, and decrease buying cycles.

An education company named X Education sells online courses to industry professionals. On any given day, many professionals who are interested in the courses land on their website and browse for courses.

The company markets its courses on several websites and search engines like Google. Once these people land on the website, they might browse the courses or fill up a form for the course or watch some videos. When these people fill up a form providing their email address or phone number, they are classified to be a lead. Moreover, the company also gets leads through past referrals. Once these leads are acquired, employees from the sales team start making calls, writing emails, etc. Through this process, some of the leads get converted while most do not.

### Business Goal

Goal from a business perspective: X Education wants to select the most promising leads, i.e. the leads that are most likely to convert into paying customers. The company requires you to build a model wherein you need to assign a lead score to each of the leads such that the customers with higher lead score have a higher conversion chance and the customers with lower lead score have a lower conversion chance. The CEO, in particular, has given a ballpark of the target lead conversion rate to be around 80%.

### Goal from a Data Scientist perspective:

Our mission is to build a better lead scoring model, targeting an 80% conversion rate and precision score. Using `predict_proba()`, we'll assess lead probabilities. This project aims to gain insights and emphasize a data-driven approach for success.

## **\*\* Prepare Work Environment\*\***

## Import Required Libraries

```
In [10]: import pandas as pd
import numpy as np
```

```
import matplotlib.pyplot as plt
import seaborn as sns
import plotly.express as px
import plotly.graph_objects as go
import warnings
from scipy.stats import linregress, uniform
from sklearn.model_selection import train_test_split, cross_val_score, StratifiedKFold
from sklearn.compose import make_column_transformer, make_column_selector
from sklearn.impute import KNNImputer, SimpleImputer
from sklearn.preprocessing import FunctionTransformer, OneHotEncoder, StandardScaler
from sklearn.pipeline import make_pipeline
from sklearn.linear_model import LogisticRegression
from sklearn.svm import SVC
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier, GradientBoostingClassifier
from sklearn.metrics import f1_score, recall_score, roc_auc_score, precision_score,
import os
os.getcwd()
```

Out[10]: 'C:\\Users\\aman'

```
In [11]: ## suppress warnings & display option
warnings.filterwarnings('ignore')
pd.set_option('display.max_columns', 50)
pd.set_option('display.max_rows', 50)
```

## 1. Load and Inspect the data

```
In [12]: df = pd.read_csv("D:\\aman_new\\Lead Scoring Assignment\\Leads.csv", encoding = 'latin1')
df.head()
```

Out[12]:

	Prospect ID	Lead Number	Lead Origin	Lead Source	Do Not Email	Do Not Call	Converted	TotalVisits	Total Time Spent on Website
0	7927b2df-8bba-4d29-b9a2-b6e0beafe620	660737	API	Olark Chat	No	No	0	0.0	
1	2a272436-5132-4136-86fa-dcc88c88f482	660728	API	Organic Search	No	No	0	5.0	67
2	8cc8c611-a219-4f35-ad23-fdfd2656bd8a	660727	Landing Page Submission	Direct Traffic	No	No	1	2.0	153
3	0cc2df48-7cf4-4e39-9de9-19797f9b38cc	660719	Landing Page Submission	Direct Traffic	No	No	0	1.0	30
4	3256f628-e534-4826-9d63-4a8b88782852	660681	Landing Page Submission	Google	No	No	1	2.0	142

1.1\* shape and info about the dataset

```
In [ ]:
In [13]: df.shape
Out[13]: (9240, 37)
In [14]: df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 9240 entries, 0 to 9239
Data columns (total 37 columns):
#   Column                                     Non-Null Count  Dtype
---  -
0   Prospect ID                               9240 non-null   object
1   Lead Number                               9240 non-null   int64
2   Lead Origin                               9240 non-null   object
3   Lead Source                               9204 non-null   object
4   Do Not Email                             9240 non-null   object
5   Do Not Call                              9240 non-null   object
6   Converted                                9240 non-null   int64
7   TotalVisits                              9103 non-null   float64
8   Total Time Spent on Website               9240 non-null   int64
9   Page Views Per Visit                      9103 non-null   float64
10  Last Activity                             9137 non-null   object
11  Country                                   6779 non-null   object
12  Specialization                            7802 non-null   object
13  How did you hear about X Education         7033 non-null   object
14  What is your current occupation            6550 non-null   object
15  What matters most to you in choosing a course 6531 non-null   object
16  Search                                    9240 non-null   object
17  Magazine                                  9240 non-null   object
18  Newspaper Article                         9240 non-null   object
19  X Education Forums                       9240 non-null   object
20  Newspaper                                 9240 non-null   object
21  Digital Advertisement                    9240 non-null   object
22  Through Recommendations                  9240 non-null   object
23  Receive More Updates About Our Courses    9240 non-null   object
24  Tags                                     5887 non-null   object
25  Lead Quality                             4473 non-null   object
26  Update me on Supply Chain Content         9240 non-null   object
27  Get updates on DM Content                 9240 non-null   object
28  Lead Profile                             6531 non-null   object
29  City                                     7820 non-null   object
30  Asymmetrique Activity Index               5022 non-null   object
31  Asymmetrique Profile Index               5022 non-null   object
32  Asymmetrique Activity Score               5022 non-null   float64
33  Asymmetrique Profile Score               5022 non-null   float64
34  I agree to pay the amount through cheque  9240 non-null   object
35  A free copy of Mastering The Interview    9240 non-null   object
36  Last Notable Activity                     9240 non-null   object
dtypes: float64(4), int64(3), object(30)
memory usage: 2.6+ MB

```

- Initial thoughts and action plan
  - check duplicates
  - drop prospect Id and lead number.
  - reformat column names without space and lowercase for practicality, and change some columns name for others names more intuitive.
  - seems that we got lot of binary columns, with only yes or no values etc. Convert it in binary encoding.

- decide what to do all the "select" data, if count it as null or assign another category like "not answered" values.

## 1.2 check if columns specified are really binary

```
In [15]: binary_cats = ['Do Not Email', 'Do Not Call', 'Search', 'Magazine', 'Newspaper Article',
                        'X Education Forums', 'Newspaper', 'Digital Advertisement', 'Through Re',
                        'Receive More Updates About Our Courses', 'Update me on Supply Chain',
                        'I agree to pay the amount through cheque', 'A free copy of Masterin']

null_values = df[binary_cats].isnull().sum()
total = df[binary_cats].count()
yes_no = df[binary_cats].applymap(lambda x: 1 if x == 'Yes' or x == 'No' else 0).sum()
df_binary_cats = pd.DataFrame({'total': total,
                              'null_%': null_values/total*100,
                              'yes/no_%': yes_no/total*100})

df_binary_cats
```

```
Out[15]:
```

	total	null_%	yes/no_%
<b>Do Not Email</b>	9240	0.0	100.0
<b>Do Not Call</b>	9240	0.0	100.0
<b>Search</b>	9240	0.0	100.0
<b>Magazine</b>	9240	0.0	100.0
<b>Newspaper Article</b>	9240	0.0	100.0
<b>X Education Forums</b>	9240	0.0	100.0
<b>Newspaper</b>	9240	0.0	100.0
<b>Digital Advertisement</b>	9240	0.0	100.0
<b>Through Recommendations</b>	9240	0.0	100.0
<b>Receive More Updates About Our Courses</b>	9240	0.0	100.0
<b>Update me on Supply Chain Content</b>	9240	0.0	100.0
<b>Get updates on DM Content</b>	9240	0.0	100.0
<b>I agree to pay the amount through cheque</b>	9240	0.0	100.0
<b>A free copy of Mastering The Interview</b>	9240	0.0	100.0

- there is no missing values
- All the columns have only "yes" or "no" values.

## 1.3. Separate train and test datasets

Let's separate train and test set before keep seeing more info. Separating train and test data is essential to avoid data leakage, evaluate model generalization, and make unbiased performance assessments in machine learning. It ensures robust model development and reliable predictions on new, unseen data.

Why stratify by target label? Stratifying train and test datasets in classification ensures balanced class representation, guarding against biased or imbalanced model learning. It promotes accurate evaluation, preventing skewed performance metrics.

```
In [16]: train, test = train_test_split(df, test_size=.2, random_state=12, stratify=df['Conv
print(f'train shape: {train.shape}')
print(f'test shape: {test.shape}')
```

```
train shape: (7392, 37)
test shape: (1848, 37)
```

## 1.4. Inspecting only training dataset

```
In [17]: print(f'In the train set are {train.duplicated().sum()} duplicates')
```

In the train set are 0 duplicates

check the values in Asymmetrique profile index columns

```
In [18]: train['Asymmetrique Profile Index'].value_counts(dropna=False)
```

```
Out[18]: Asymmetrique Profile Index
NaN      3362
02.Medium 2243
01.High   1762
03.Low     25
Name: count, dtype: int64
```

```
In [19]: train['Asymmetrique Activity Index'].value_counts(dropna=False)
```

```
Out[19]: Asymmetrique Activity Index
NaN      3362
02.Medium 3080
01.High   648
03.Low    302
Name: count, dtype: int64
```

Assymetrique's columns treatment:

We have identified three distinct categories and some missing records. To improve the data's representativeness for machine learning modeling we will focus on the integer values and reverse their order, emphasizing a higher-is-better perspective.

- 

High: Assigned a value of \* 3 Medium: Assigned a value o \*f 2 Low: Assigned a value of 1

## 2. Data cleaning & Feature Engineering

### 2.1 Data Cleaning

Let us embark on our first data cleaning endeavor! Our strategy involves transforming each step into Scikit-learn transformation objects, harmonizing the entire process into a unified pipeline.

```
In [20]: def data_cleaning(df):
  """Do some of the data cleaning procedures that we
  specified at the begining of the notebook"""
  # drop columns id columns
  df = df.drop(['Prospect ID', 'Lead Number'], axis=1)

  # asymmetrique index columns transformation
  df['Asymmetrique Activity Index'] = df['Asymmetrique Activity Index'].str.split('.')
  .str.replac
  .str.replac
  .astype(np.
  )
  df['Asymmetrique Profile Index'] = df['Asymmetrique Profile Index'].str.split('.')
  .str.replac
  .str.replac
  .astype(np.
  )

  # binary encoding
  df[binary_cats] = df[binary_cats].applymap(lambda x: 0 if x == 'No' else 1)

  # rename columns for praticicity
  df.columns = df.columns.str.replace(' ', '_').str.lower()
  return df

# Convert custom function into transformer
initial_clean = FunctionTransformer(data_cleaning)

train_clean = initial_clean.fit_transform(train);
```

### 4.2 Inspecting category columns

In this stage, we'll first inspect the categorical columns from a practical and business-oriented perspective, before delving into more advanced statistical analysis.

I firmly believe that simplicity often holds the key to effective solutions.

The goal is to take a first look through all category columns to do some feature engineering, extract some initial thoughts for future EDA/feature engineer and handling missing values.

```
In [21]: train_clean.lead_origin.value_counts(dropna=False)
```

```
Out[21]: lead_origin
Landing Page Submission    3906
API                        2889
Lead Add Form              550
Lead Import                47
Name: count, dtype: int64
```

```
In [22]: train_clean.lead_source.value_counts(dropna=False)
```

```
Out[22]: lead_source
Google                    2326
Direct Traffic            2033
Olark Chat                1408
Organic Search            916
Reference                 405
Welingak Website         111
Referral Sites            98
Facebook                  46
NaN                       27
bing                      6
Click2call                4
google                    3
Live Chat                 2
blog                      1
testone                   1
Social Media              1
youtubechannel            1
WeLearn                   1
Press_Release             1
Pay per Click Ads         1
Name: count, dtype: int64
```

```
In [23]: train_clean.last_activity.value_counts(dropna=False)
```

```
Out[23]: last_activity
Email Opened              2712
SMS Sent                  2224
Olark Chat Conversation   793
Page Visited on Website  506
Converted to Lead         336
Email Bounced            271
Email Link Clicked        211
Form Submitted on Website 97
NaN                       80
Unreachable              71
Unsubscribed              50
Had a Phone Conversation  23
Approached upfront        8
View in browser link Clicked 5
Email Received            2
Resubscribed to emails    1
Email Marked Spam         1
Visited Booth in Tradeshow 1
Name: count, dtype: int64
```



```
In [24]: train_clean.country.value_counts(dropna=False)
```

```
Out[24]: country
India          5201
NaN            1954
United States    57
United Arab Emirates  49
Singapore        21
Saudi Arabia     17
United Kingdom   13
Australia         9
Qatar             9
Bahrain           6
Oman              4
Nigeria           4
Germany           4
France            4
unknown          4
Kuwait            3
Hong Kong         3
Canada            3
South Africa      3
China             2
Belgium           2
Sweden            2
Italy             2
Bangladesh        2
Ghana             2
Netherlands       2
Switzerland       1
Denmark           1
Uganda            1
Tanzania          1
Russia            1
Philippines       1
Malaysia          1
Sri Lanka         1
Asia/Pacific Region 1
Kenya             1
Name: count, dtype: int64
```

```
In [25]: train_clean.specialization.value_counts(dropna=False)
```

```
Out[25]: specialization
Select      1554
NaN         1154
Finance Management    780
Marketing Management  682
Human Resource Management  661
Operations Management  407
Business Administration  329
IT Projects Management  295
Supply Chain Management  281
Banking, Investment And Insurance  272
Media and Advertising  161
Travel and Tourism     155
International Business  152
Healthcare Management  128
Hospitality Management  89
E-COMMERCE           84
Retail Management     76
Rural and Agribusiness  59
E-Business            43
Services Excellence   30
Name: count, dtype: int64
```

```
In [26]: train_clean.how_did_you_hear_about_x_education.value_counts(dropna=False)
```

```
Out[26]: how_did_you_hear_about_x_education
Select      4023
NaN         1768
Online Search    647
Word Of Mouth    288
Student of SomeSchool  258
Other           139
Multiple Sources  128
Social Media     51
Advertisements   49
SMS              21
Email            20
Name: count, dtype: int64
```

```
In [27]: train_clean.what_is_your_current_occupation.value_counts(dropna=False)
```

```
Out[27]: what_is_your_current_occupation
Unemployed      4474
NaN             2159
Working Professional  559
Student          172
Other            14
Housewife         8
Businessman        6
Name: count, dtype: int64
```

```
In [28]: train_clean.what_matters_most_to_you_in_choosing_a_course.value_counts(dropna=False)
```

```
Out[28]: what_matters_most_to_you_in_choosing_a_course
Better Career Prospects      5216
NaN                          2173
Flexibility & Convenience      2
Other                          1
Name: count, dtype: int64
```

```
In [29]: train_clean.tags.value_counts(dropna=False)
```

```
Out[29]: tags
NaN                                2687
Will revert after reading the email 1664
Ringing                            958
Interested in other courses         418
Already a student                   365
Closed by Horizzon                  287
switched off                        200
Busy                               146
Lost to EINS                        138
Not doing further education         110
Interested in full time MBA          92
Graduation in progress              87
invalid number                      60
Diploma holder (Not Eligible)        51
wrong number given                  42
opp hangup                          29
number not provided                 21
in touch with EINS                   7
Want to take admission but has financial problems 6
Lost to Others                      5
Still Thinking                      4
In confusion whether part time or DLP 4
Interested in Next batch             3
Lateral student                     3
University not recognized            2
Shall take in the next coming month  2
Recognition issue (DEC approval)    1
Name: count, dtype: int64
```

```
In [30]: train_clean.lead_quality.value_counts(dropna=False)
```

```
Out[30]: lead_quality
NaN                3796
Might be           1249
Not Sure           884
High in Relevance  508
Worst              490
Low in Relevance   465
Name: count, dtype: int64
```

```
In [31]: train_clean.city.value_counts(dropna=False)
```

```
Out[31]: city
Mumbai                2581
Select                1806
NaN                   1139
Thane & Outskirts      607
Other Cities           539
Other Cities of Maharashtra 349
Other Metro Cities     307
Tier II Cities         64
Name: count, dtype: int64
```

```
In [32]: train_clean.last_notable_activity.value_counts(dropna=False)
```

```
Out[32]: last_notable_activity
Modified                2734
Email Opened           2219
SMS Sent               1759
Page Visited on Website 259
Olark Chat Conversation 153
Email Link Clicked     140
Email Bounced          49
Unsubscribed           39
Unreachable            24
Had a Phone Conversation 10
Form Submitted on Website 1
Email Received          1
View in browser link Clicked 1
Resubscribed to emails  1
Email Marked Spam       1
Approached upfront      1
Name: count, dtype: int64
```

## Initial feature Engineering

Apply initial changes described in the previous insights

```
In [33]: def initial_feature_engineering(df):
        """Do some feature engineering"""
        # Lead_source
        df['lead_source'] = df['lead_source'].str.replace('|'.join(['google','Pay per Cli
        df['lead_source'] = df['lead_source'].apply(lambda x: "Referral Sites" if 'blog'
        df['lead_source'] = df['lead_source'].str.replace('Live Chat','Olark Chat')
        df['lead_source'] = df['lead_source'].str.replace('bing','Organic Search')
        df['lead_source'] = df[df['lead_source'] != 'Other'].lead_source.apply(lambda x:
        # Last_activity and Last_notable_activity
        activity = ['last_activity','last_notable_activity']
        df[activity] = df[activity].apply(lambda x: x.str.replace('|'.join(['Email Receiv
        df[activity] = df[activity].apply(lambda x: x.str.replace('|'.join(['Email Marked
        df[activity] = df[activity].apply(lambda x: x.str.replace('Resubscribed to emails
        df[activity] = df[activity].apply(lambda x: x.str.replace('|'.join(['Visited Boot
        # country
        df['country'] = df['country'].apply(lambda x: np.nan if x in ['Unknown','unknown'
        # specialization
        df['specialization'] = df['specialization'].str.replace('|'.join(['E-COMMERCE','E
```

```

df['specialization'] = df['specialization'].str.replace('Banking, Investment And
df['specialization'] = df['specialization'].str.replace('Media and Advertising', '
df['specialization'] = df['specialization'].str.replace('Select', 'Not Provided')
# how_did_you_hear
df['how_did_you_hear_about_x_education'] = df['how_did_you_hear_about_x_education
df['how_did_you_hear_about_x_education'] = df['how_did_you_hear_about_x_education
# importance_in_course
df['what_matters_most_to_you_in_choosing_a_course'] = df['what_matters_most_to_yo
# Lead_profile
df['lead_profile'] = df['lead_profile'].str.replace('Select', 'Not Assigned')
# city
df['city'] = df['city'].str.replace('Select', 'Not Provided')

return df

initial_feature_engineering = FunctionTransformer(initial_feature_engineering)
train_clean = initial_feature_engineering.fit_transform(train_clean)

```

### 3. Explore Missing values

copy of the dataset and visualizations style

```

In [34]: train_ = train_clean.copy()

# Set style for better visualizations
train_eda = train_clean.copy()
sns.set_style('dark')
sns.set(rc={'axes.grid': False})
sns.set_palette('viridis')

```

```

In [35]: null_ = pd.DataFrame()
null_['proportion'] = np.round(train_clean.isnull().sum()/len(train_clean),4) * 100
null_['amount'] = train_clean.isnull().sum()

# Show only those columns with at least 1 missing value
null_.sort_values(by='proportion', ascending=False)[null_.amount > 0]

```

Out[35]:

	proportion	amount
lead_quality	51.35	3796
asymmetrique_activity_index	45.48	3362
asymmetrique_profile_score	45.48	3362
asymmetrique_profile_index	45.48	3362
asymmetrique_activity_score	45.48	3362
tags	36.35	2687
lead_profile	29.40	2173
what_matters_most_to_you_in_choosing_a_course	29.40	2173
what_is_your_current_occupation	29.21	2159
country	26.50	1959
how_did_you_hear_about_x_education	23.92	1768
specialization	15.61	1154
city	15.41	1139
page_views_per_visit	1.45	107
totalvisits	1.45	107
last_activity	1.08	80

\*\* insights \* missing values in certain columns, often requiring employee input, might stem from uncategorized leads. streamlining lead management can improve data collection, inform decision-making, and optimize lead conversion strategies. Further investigation is necessary to confirm this hypothesis

## define plot functions

```
In [36]: def barplot_catcols(column,width,heigh):
    """Plot conversion rate"""
    fig, ax = plt.subplots(figsize=(width,heigh))
    ax = sns.barplot(data=train_.fillna('NaN'), x='converted', y=column,
                     order=order(train_.fillna('NaN'),column),
                     orient='h', palette='viridis',
                     seed=2)
    plt.title(f'Conversion Rate by {column.replace("_"," ").title()}', loc='left', si
    return ax

def order(df,x,y=None):
    if y is not None:
        return df.groupby(x)[y].mean().sort_values(ascending=False).index
```

```

else:
    return df.groupby(x)['converted'].mean().sort_values(ascending=False).index

```


## How much of the missing values belong to the same people?

```

In [37]: # Number of missing values in each row
train_['amount_missing'] = train_.isnull().sum(1)

# Plot the relation between amount missing and conversion rate
fig, ax = plt.subplots(figsize=(8,5))
ax = sns.barplot(data=train_.fillna('NaN'), x='converted', y='amount_missing',
                 orient='h', palette='viridis',
                 seed=2)
plt.title(f'Conversion Rate by Amount Missing', loc='left', size=20)
plt.show()

```

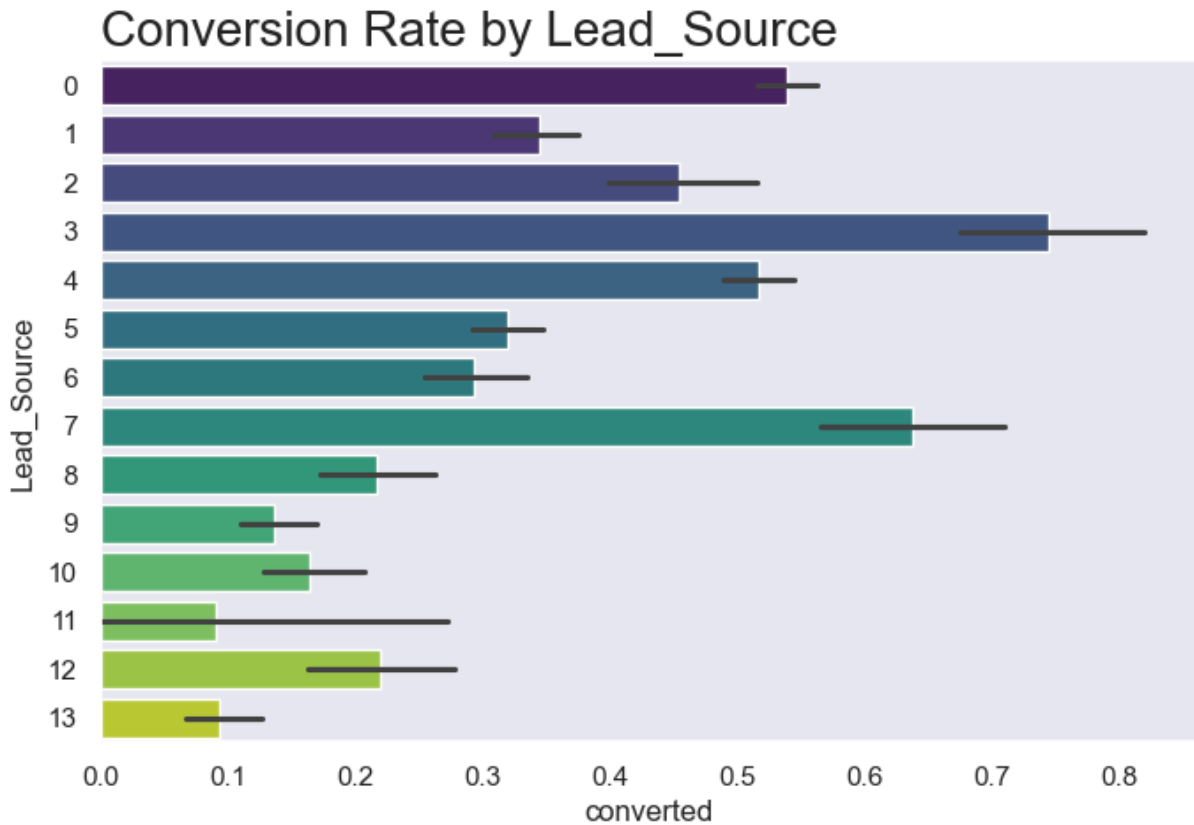


```


In [38]: # Number of missing values in each row
train_['Lead_Source'] = train_.isnull().sum(1)

# Plot the relation between amount missing and conversion rate
fig, ax = plt.subplots(figsize=(8,5))
ax = sns.barplot(data=train_.fillna('NaN'), x='converted', y='Lead_Source',
                 orient='h', palette='viridis',
                 seed=2)
plt.title(f'Conversion Rate by Lead_Source', loc='left', size=20)
plt.show()

```



```
In [39]: fig, ax = plt.subplots(figsize=(8,2))
ax = sns.barplot(data=train_, x='amount_missing', y='converted',
                orient='h', palette=sns.color_palette('viridis',2),
                seed=2)
plt.title(f'Amount missing by leads conversion', loc='left', size=18)
plt.show()
```



## 3.2 Correlation of numerical columns with converted column¶

```
In [40]: correlations = train_.select_dtypes('number').corr()['converted'].sort_values(ascen

plt.figure(figsize=(8, 8))
correlations[1:].plot(kind='barh',
```

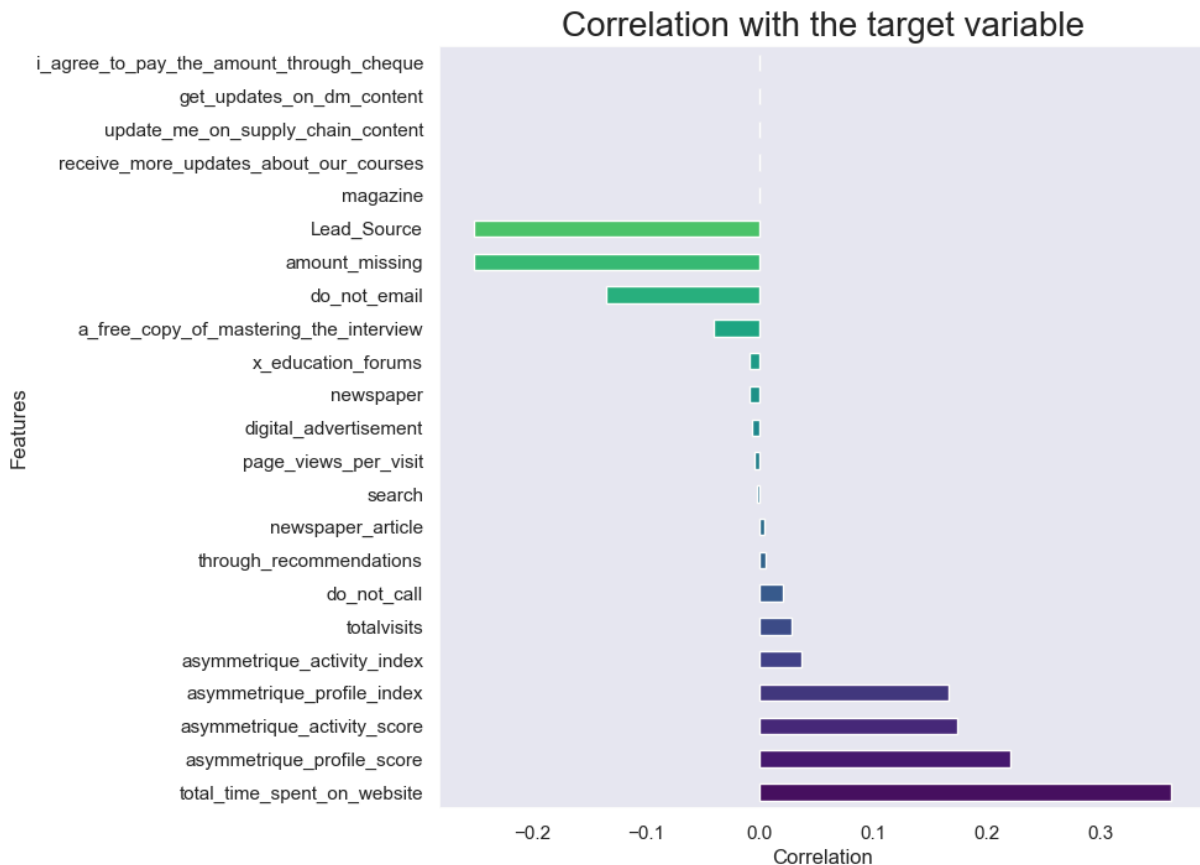


```

color=sns.color_palette('viridis', len(correlations)))

plt.title('Correlation with the target variable', fontsize=20)
plt.xlabel('Correlation')
plt.ylabel('Features')
plt.show()

```



```

In [41]: print(f'Duplicate rows from original dataset: {train.duplicated().sum()}')
print(f'Duplicate rows after feature engineer: {train_clean.duplicated().sum()}')

```

Duplicate rows from original dataset: 0

Duplicate rows after feature engineer: 984

Handle Missing Values: We currently lack sufficient information to determine the best approach for dealing with missing values. To address this, we will conduct a detailed data exploration, searching for patterns related to lead conversion. Once we have a clearer understanding, we can devise the most appropriate strategy for handling these missing records.

## 4. Exploratory Data Analysis


¶ Considering the prevalence of categorical or binary variables, we'll treat "NaN" values as a distinct category for comparison. For numerical columns with few "NaN" values, we'll exclude them to ensure robust analysis. This follows EDA best practices for gaining valuable insights from the dataset.

```
In [42]: count = train_['converted'].value_counts()

fig, ax = plt.subplots(figsize=(10, 5))
ax.pie(count, labels=count.index, autopct='%1.1f%%', startangle=90, colors=['#29568',
ax.set_title('Converted', size=20)

centre_circle = plt.Circle((0,0),0.70,fc='white')
fig.gca().add_artist(centre_circle)

plt.axis('equal')
plt.show()
```



Insight: The dataset exhibits a relatively balanced distribution of the target variable. While there may be some variations in class proportions, it's not extremely unbalanced.

```
In [43]: train_.loc[:, 'asymmetrique_activity_index': 'asymmetrique_profile_score'].corr().sty
```

Out[43]:

	asymmetrique_activity_index	asymmetrique_profile_index	asyn
asymmetrique_activity_index	1.000000	-0.145399	
asymmetrique_profile_index	-0.145399	1.000000	
asymmetrique_activity_score	0.855985	-0.145366	
asymmetrique_profile_score	-0.122669	0.883177	

Insight: As expected, there's a strong correlation between the "Score" and "Index" columns. Given the level of detail in the data, retaining the score columns appears to be a sound choice. These columns appear to offer valuable information, and their inclusion in our analysis is likely to yield valuable insights.


## 4.1 . Categorical Variables

```
In [44]: fig, ax = plt.subplots(1,2, figsize=(12,6), sharey=True)

sns.barplot(data=train_.fillna('NaN'), x='lead_profile', y='converted',
            palette='viridis', order=order(train_.fillna('NaN'),'lead_profile'),
            seed=2, ax=ax[0])
ax[0].set_xticklabels(ax[0].get_xticklabels(), rotation=90)
ax[0].set_title(f'Conversion Rate by Lead Profile', loc='left', size=16)

sns.barplot(data=train_.fillna('NaN'), x='asymmetrique_profile_score', y='converted',
            palette='viridis', order=order(train_.fillna('NaN'),'asymmetrique',
            seed=2, ax=ax[1])
ax[1].set_xticklabels(ax[1].get_xticklabels(), rotation=90)
ax[1].set_title(f'Conversion Rate by Asymmetrique Profile Score', loc='left', size=

plt.tight_layout()
plt.show()
```



Insights: There's a significant difference in the conversion rate of people with "Not Assigned" and "NaN" values, which might suggest that they do not belong to the same group. Profile Score could be a better predictor than Lead Profile, as the conversion rate tends to increase with higher scores. Because both columns essentially represent the same information, it is advisable to drop the Lead Profile column for simplicity and clarity in the analysis.

## 4.2 Lead Activity

- Correlation between activity track record (columns related with the web) and activity/profile score

```
In [45]: activity_columns = ['totalvisits', 'total_time_spent_on_website', 'page_views_per_visit',
                             'asymmetrique_profile_score', 'asymmetrique_activity_score']

train_[activity_columns].corr().style.background_gradient(cmap='vlag_r')
```

Out[45]:

	totalvisits	total_time_spent_on_website	page_views_per_visit	asymmetrique_profile_score	asymmetrique_activity_score
totalvisits	1.000000	0.261952	0.598883	0.129016	-0.061397
total_time_spent_on_website	0.261952	1.000000	0.323684	0.167992	-0.066008
page_views_per_visit	0.598883	0.323684	1.000000	0.165945	-0.171264
asymmetrique_profile_score	0.129016	0.167992	0.165945	1.000000	
asymmetrique_activity_score	-0.061397	-0.066008	-0.171264		1.000000

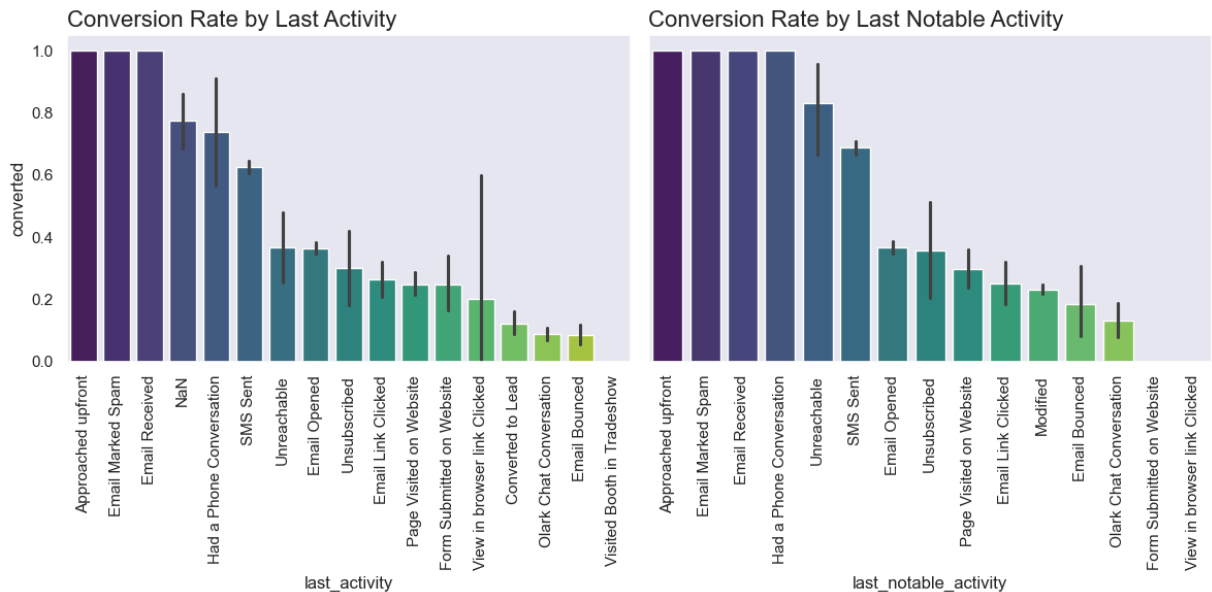
- Having columns about last activity and last notable activity provides more information?

```
In [46]: fig, ax = plt.subplots(1,2, figsize=(12,6), sharey=True)

sns.barplot(data=train_.fillna('NaN'), x='last_activity', y='converted',
             order=order(train_.fillna('NaN'),'last_activity'),
             palette='viridis',
             seed=2, ax=ax[0])
ax[0].set_xticklabels(ax[0].get_xticklabels(), rotation=90)
ax[0].set_title(f'Conversion Rate by Last Activity', loc='left', size=16)

sns.barplot(data=train_.fillna('NaN'), x='last_notable_activity', y='converted',
             order=order(train_.fillna('NaN'),'last_notable_activity'),
             palette='viridis', seed=2)
ax[1].set_xticklabels(ax[1].get_xticklabels(), rotation=90)
ax[1].set_title(f'Conversion Rate by Last Notable Activity', loc='left', size=16)

plt.tight_layout()
plt.show()
```



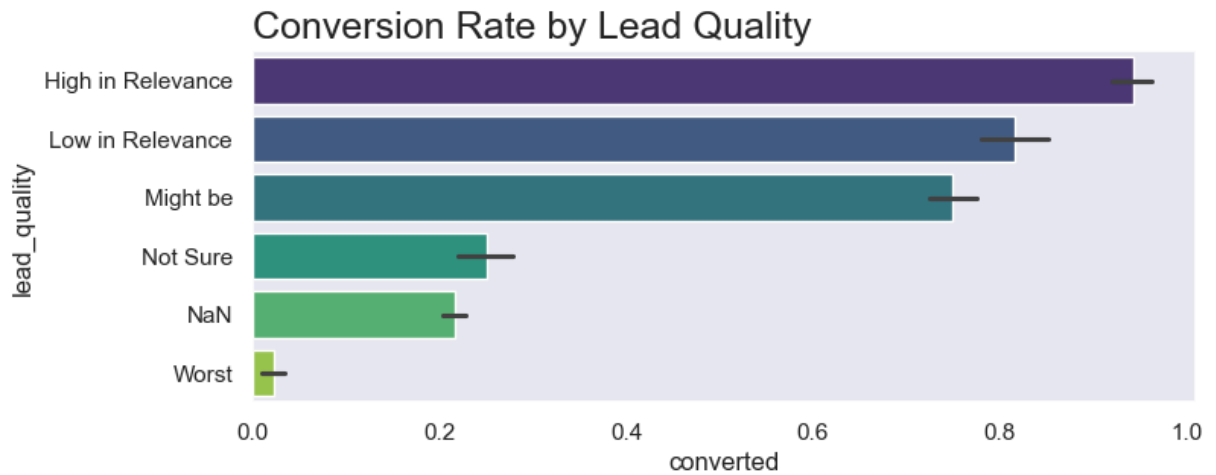
Insights: Activity Score seems to be a less effective predictor compared to Profile Score. No significant relationship was found between Activity Index and Last Activity, Last Notable Activity, or columns related to visits. There's no significant correlation among columns related to visits. Last Activity does not appear to provide substantially more information than Last Notable Activity. Hence, it may be preferable to retain Last Notable Activity.

## \*\* Business Suggestion:

\*Our analysis reveals a significant correlation between phone conversations and lead conversions. To maximize results, consider increasing phone calls to leads. Prioritizing "Hot Leads" for calls can enhance resource allocation and boost conversion rates, ultimately driving better business outcomes.

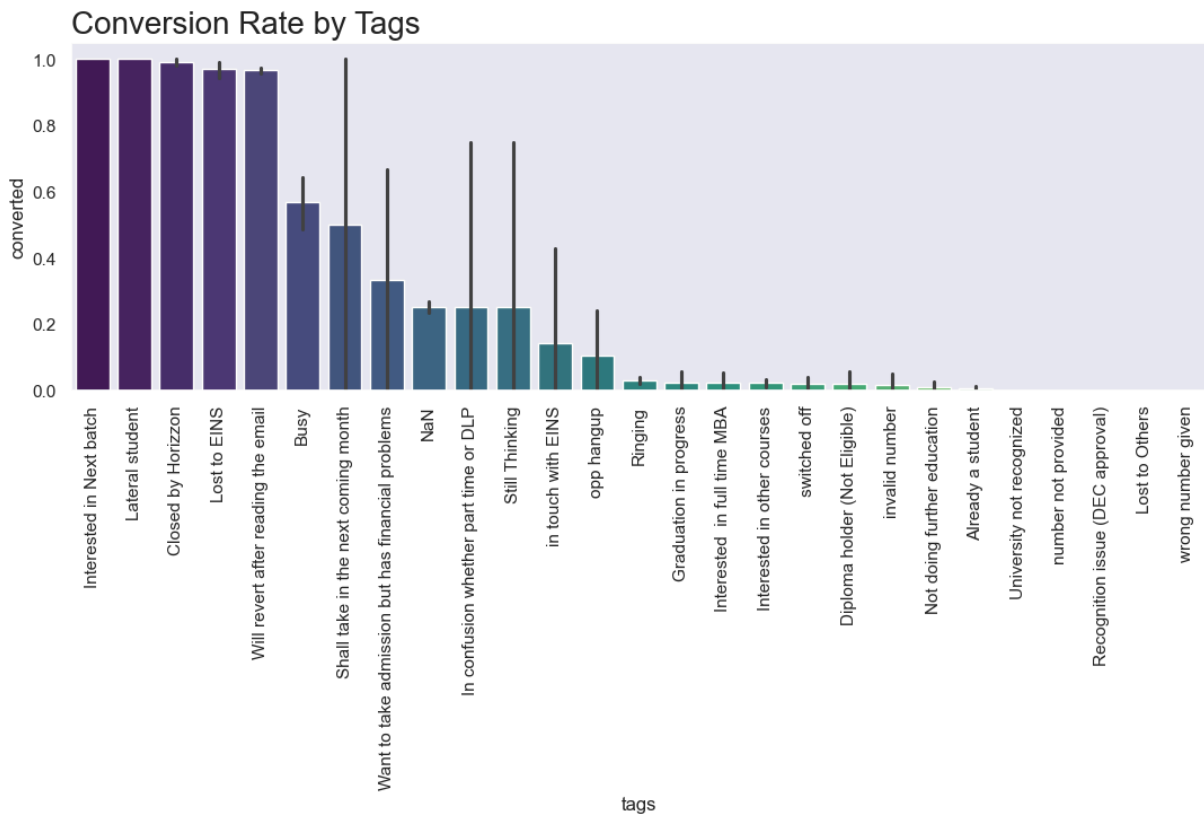
## 4.3 Lead Quality and Tags

```
In [47]: barplot_catcols('lead_quality',8,3)
plt.show()
```



```
In [48]: fig, ax = plt.subplots(figsize=(13,4))

sns.barplot(data=train_.fillna('NaN'), x='tags', y='converted',
            order=order(train_.fillna('NaN'),'tags'),
            palette='viridis',
            seed=2)
plt.xticks(rotation=90)
plt.title(f'Conversion Rate by Tags', loc='left', size=20)
plt.show()
```



## 4.4 Occupation and Specialization

```
In [49]: fig, ax = plt.subplots(1,2, figsize=(14,7), sharey=True)
```

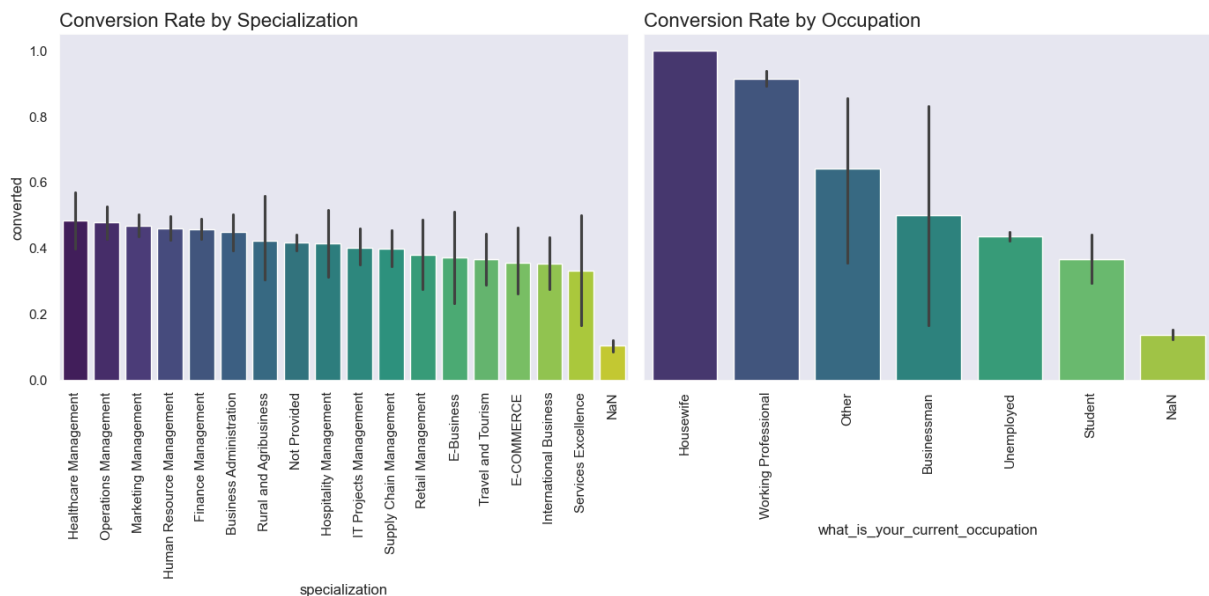
```

sns.barplot(data=train_.fillna('NaN'), x='specialization', y='converted',
            order=order(train_.fillna('NaN'),'specialization'),
            palette='viridis',
            seed=2, ax=ax[0])
ax[0].set_xticklabels(ax[0].get_xticklabels(), rotation=90)
ax[0].set_title(f'Conversion Rate by Specialization', loc='left', size=16)

sns.barplot(data=train_.fillna('NaN'), x='what_is_your_current_occupation', y='conv
            order=order(train_.fillna('NaN'),'what_is_your_current_occupation
            palette='viridis', seed=2)
ax[1].set_xticklabels(ax[1].get_xticklabels(), rotation=90)
ax[1].set_title(f'Conversion Rate by Occupation', loc='left', size=16)

plt.tight_layout()
plt.show()

```



- Number of missing values for each row in these two categories

```
In [50]: train[['what_is_your_current_occupation', 'specialization']].isnull().sum(1).value_
```

```
Out[50]: 0    5220
         2    1141
         1    1031
         Name: count, dtype: int64
```

## \*\* 4.5. The source from which the customer heard about X Education and the source of the lead\*\*

```

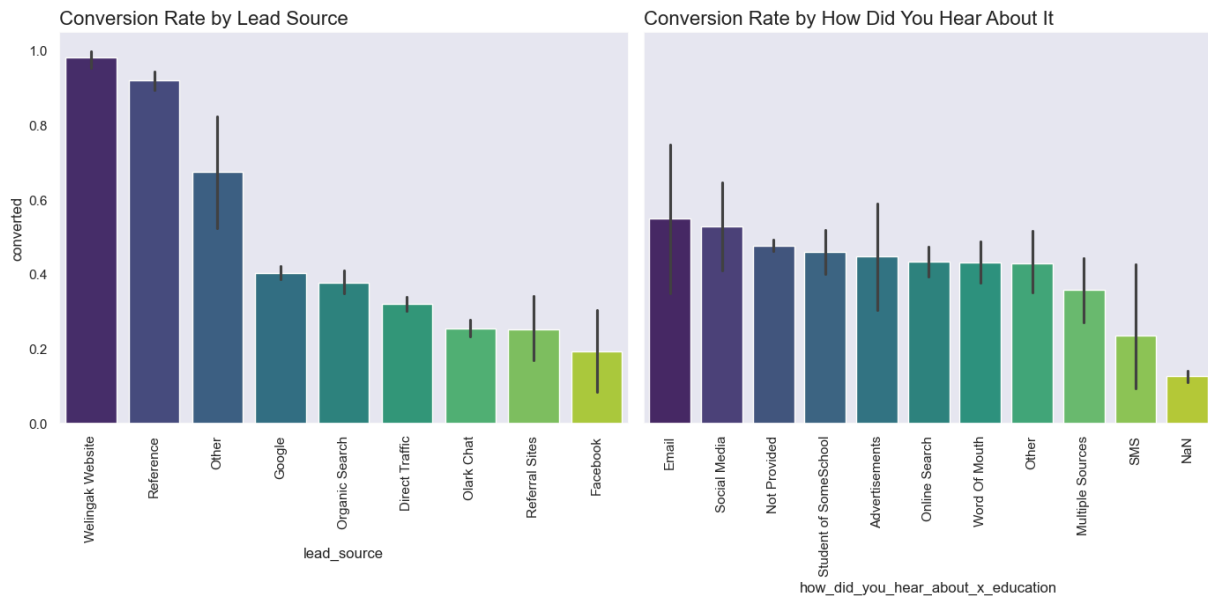
In [51]: fig, ax = plt.subplots(1,2, figsize=(14,7), sharey=True)

sns.barplot(data=train_.fillna('NaN'), x='lead_source', y='converted',
            order=order(train_.fillna('NaN'),'lead_source'),
            palette='viridis',
            seed=2, ax=ax[0])
ax[0].set_xticklabels(ax[0].get_xticklabels(), rotation=90)
ax[0].set_title(f'Conversion Rate by Lead Source', loc='left', size=16)

```

```
sns.barplot(data=train_.fillna('NaN'), x='how_did_you_hear_about_x_education', y='converted',
            order=order(train_.fillna('NaN'),'how_did_you_hear_about_x_education'),
            palette='viridis', seed=2)
ax[1].set_xticklabels(ax[1].get_xticklabels(), rotation=90)
ax[1].set_title(f'Conversion Rate by How Did You Hear About It', loc='left', size=14)

plt.tight_layout()
plt.show()
```



## Business Suggestion:

Referrals, with a 90% conversion rate, are a top-performing lead source due to their trustworthiness. To capitalize on this potential, the business should incentivize, personalize, track, showcase testimonials, and leverage word-of-mouth marketing for effective growth.

## Numeric Variables

```
In [52]: train_.select_dtypes(include=['number']).unique().sort_values()
```



```

Out[52]: i_agree_to_pay_the_amount_through_cheque      1
         get_updates_on_dm_content                    1
         update_me_on_supply_chain_content             1
         receive_more_updates_about_our_courses        1
         magazine                                       1
         do_not_email                                  2
         a_free_copy_of_mastering_the_interview        2
         through_recommendations                      2
         newspaper                                      2
         digital_advertisement                         2
         newspaper_article                             2
         search                                         2
         converted                                      2
         do_not_call                                   2
         x_education_forums                            2
         asymmetrique_activity_index                   3
         asymmetrique_profile_index                   3
         asymmetrique_profile_score                   10
         asymmetrique_activity_score                   11
         amount_missing                                14
         Lead_Source                                   14
         totalvisits                                    40
         page_views_per_visit                           103
         total_time_spent_on_website                   1635
         dtype: int64

```

## columns Related to Web Visits

```

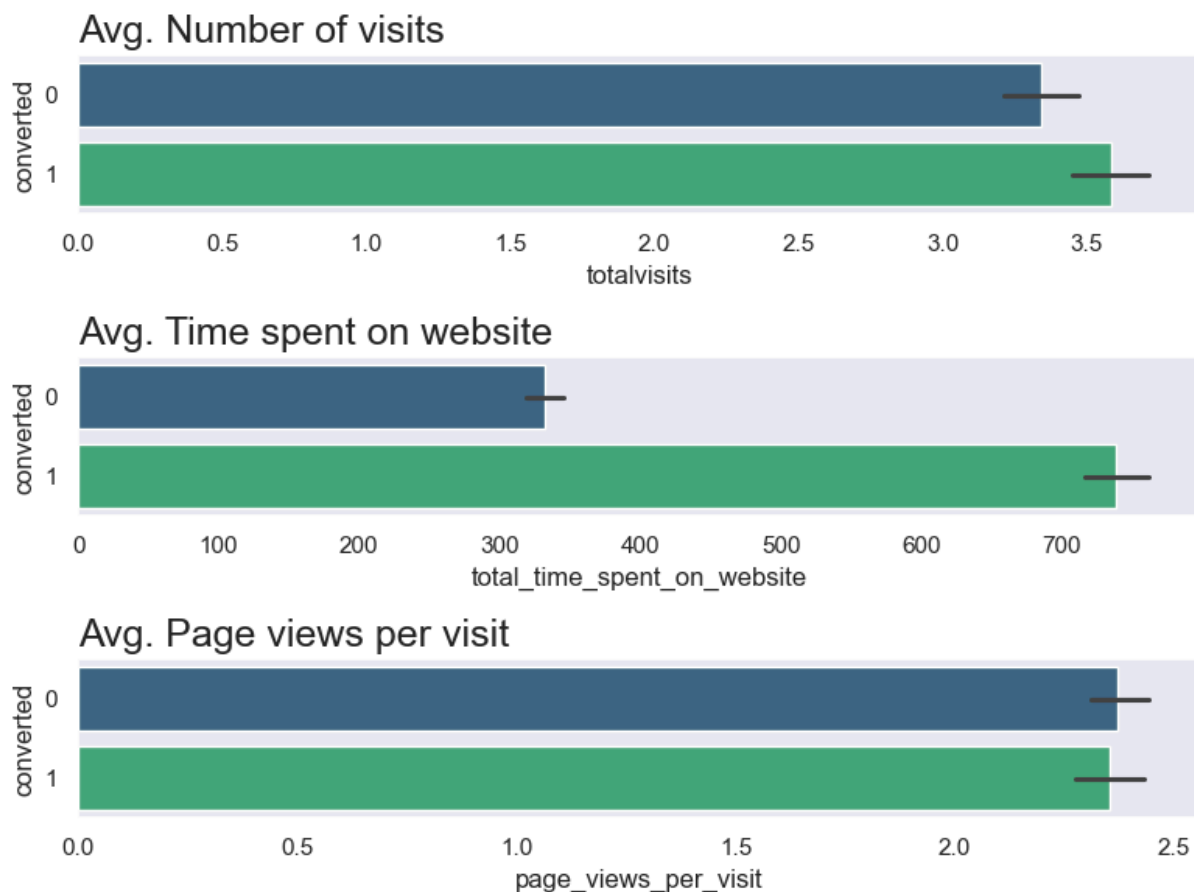
In [53]: fig, ax = plt.subplots(3, figsize=(8,6))
         sns.barplot(data=train_, x='totalvisits', y='converted',
                     orient='h', palette='viridis',
                     seed=2, ax=ax[0])
         ax[0].set_title(f'Avg. Number of visits', loc='left', size=18)

         sns.barplot(data=train_, x='total_time_spent_on_website', y='converted',
                     orient='h', palette='viridis',
                     seed=2, ax=ax[1])
         ax[1].set_title(f'Avg. Time spent on website', loc='left', size=18)

         sns.barplot(data=train_, x='page_views_per_visit', y='converted',
                     orient='h', palette='viridis',
                     seed=2, ax=ax[2])
         ax[2].set_title(f'Avg. Page views per visit', loc='left', size=18)

         plt.tight_layout()
         plt.show()

```




```
In [54]: fig, ax = plt.subplots(3,1, figsize=(8,6))
sns.boxplot(data=train_, x='totalvisits',
            ax=ax[0], palette='viridis')
ax[0].set_title('Total Visits', loc='left', size=16)

sns.boxplot(data=train_, x='total_time_spent_on_website',
            ax=ax[1], palette='viridis')
ax[1].set_title('Time spent on web', loc='left', size=16)

sns.boxplot(data=train_, x='page_views_per_visit',
            ax=ax[2], palette='viridis')
ax[2].set_title('Page views per visit', loc='left', size=16)

plt.tight_layout()
plt.show()
```


Total Visits



Time spent on web



Page views per visit



## Insights:

There's a significant difference in conversion rate among both groups. Leads that convert more spent much more time on the website

## 5. Data Wrangling

### Outlier treatments

Addressing outliers in TotalVisits and Page Views Per Visit is essential for model performance, particularly in Logistic Regression. Capping these variables at the 95th percentile is recommended for model stability and preventing inflated coefficients. It enhances model generalization in various classification models like Decision Trees, Random Forests, and Support Vector Machines.

### Missing Values Strategies

- **Numeric Columns (KNN Imputation):** Utilizing KNNImputer for imputing missing values in Total Visits and Page Views Per Visit is a preferable choice over median, mean, or mode imputation. KNNImputer considers feature relationships, preserving data distribution, and handling multicollinearity effectively.

- Categorical Columns (Missing Category): Treating missing values as a separate category, rather than imputing with the mode, maintains data integrity, avoids biases, and improves model reliability and accuracy, especially considering the significant difference in conversion rate between leads with missing records and others.

## 5.1 Feature Engineer

apply all the insights discovered during EDA.

```
In [55]: def eda_feature_engineering(df):
# tags column
df['tags'] = df['tags'].str.replace('|'.join(['invalid number','wrong number give
df['tags'] = df['tags'].str.replace('|'.join(["In confusion whether part time or
df['tags'] = df['tags'].str.replace("University not recognized","Not eligible")
df['tags'] = df[df['tags'].notnull()].tags.apply(lambda x: 'Not eligible' if 'hol
df['tags'] = df['tags'].str.replace('|'.join(["Interested in other courses", "Int
df['tags'] = df['tags'].str.replace('|'.join(["Ringing","switched off"]), "Still n
df['tags'] = df['tags'].str.replace('|'.join(["Want to take admission but has fin
df['tags'] = df[df['tags'].notnull()].tags.apply(lambda x: 'Not eligible for the
df['tags'] = df[df['tags'].notnull()].tags.apply(lambda x: 'Other' if x not in df

# country and city
indian_cities = ['Mumbai','Thane & Outskirts','Other Cities of Maharashtra','Tier
df.loc[(df.country != 'India') & (df.city.isin(indian_cities)), 'country'] = 'Indi
df['country'] = df.loc[df['country'].notnull(), 'country'].apply(lambda x: 'Other'

# Lead quality
df['lead_quality'] = df['lead_quality'].fillna('Not Sure')

# convert asymmetrique index columns in strings columns
df[['asymmetrique_profile_index', 'asymmetrique_activity_index']] = df[['asymmetri

# drop columns with unique values
drop_cols = ['magazine', 'receive_more_updates_about_our_courses', 'update_me_on_su
            'get_updates_on_dm_content', 'i_agree_to_pay_the_amount_through_chequ
df = df.drop(drop_cols, axis=1)

#add amount_missing column
df['amount_missing'] = df.isnull().sum(1)
return df

eda_feature_engineering = FunctionTransformer(eda_feature_engineering)
```

## 5.2 Handling Outliers

```
In [56]: def cap_outliers(df):
        """Replace outliers with the 95th percentile"""
        num_cols = ['totalvisits', 'page_views_per_visit', 'total_time_spent_on_website']
        df[num_cols[0]].apply(lambda x: df[num_cols[0]].quantile(.95) if x > df[num_cols[0]
        df[num_cols[1]].apply(lambda x: df[num_cols[1]].quantile(.95) if x > df[num_cols[1]
        df[num_cols[2]].apply(lambda x: df[num_cols[2]].quantile(.95) if x > df[num_cols[2]
        return df
```

```
cap_outliers = FunctionTransformer(cap_outliers);
```

### 5.3 Handling missing values and scaling columns for modeling

- 1. Apply OneHotEncoder to all the categorical columns.
- 2.

Apply StandardScaler to the numeric columns if there aren't binary.

```
In [57]: cat_columns = ['lead_origin', 'lead_source', 'country', 'what_is_your_current_occupati
            'what_matters_most_to_you_in_choosing_a_course', 'tags', 'lead_qualit
            'city', 'last_notable_activity']

num_cols = ['totalvisits', 'page_views_per_visit', 'total_time_spent_on_website',
            'asymmetrique_activity_score', 'asymmetrique_profile_score', 'amount_miss

impute_knn = KNNImputer(n_neighbors=5)
impute_cons = SimpleImputer(strategy='constant', fill_value='Missing')
ohe = OneHotEncoder(handle_unknown='ignore')
sc = StandardScaler()

# Make pipelines for both type of columns treatments
pipe_cat = make_pipeline(impute_cons, ohe)
pipe_num = make_pipeline(sc, impute_knn)

impute_scale = make_column_transformer(
    (pipe_cat, cat_columns),
    (pipe_num, num_cols),
    remainder='drop'
)
```

### 5.4 Saperate X and Y

```
In [58]: X_train = train.drop('Converted', axis=1)
y_train = train.loc[:, 'Converted']
```

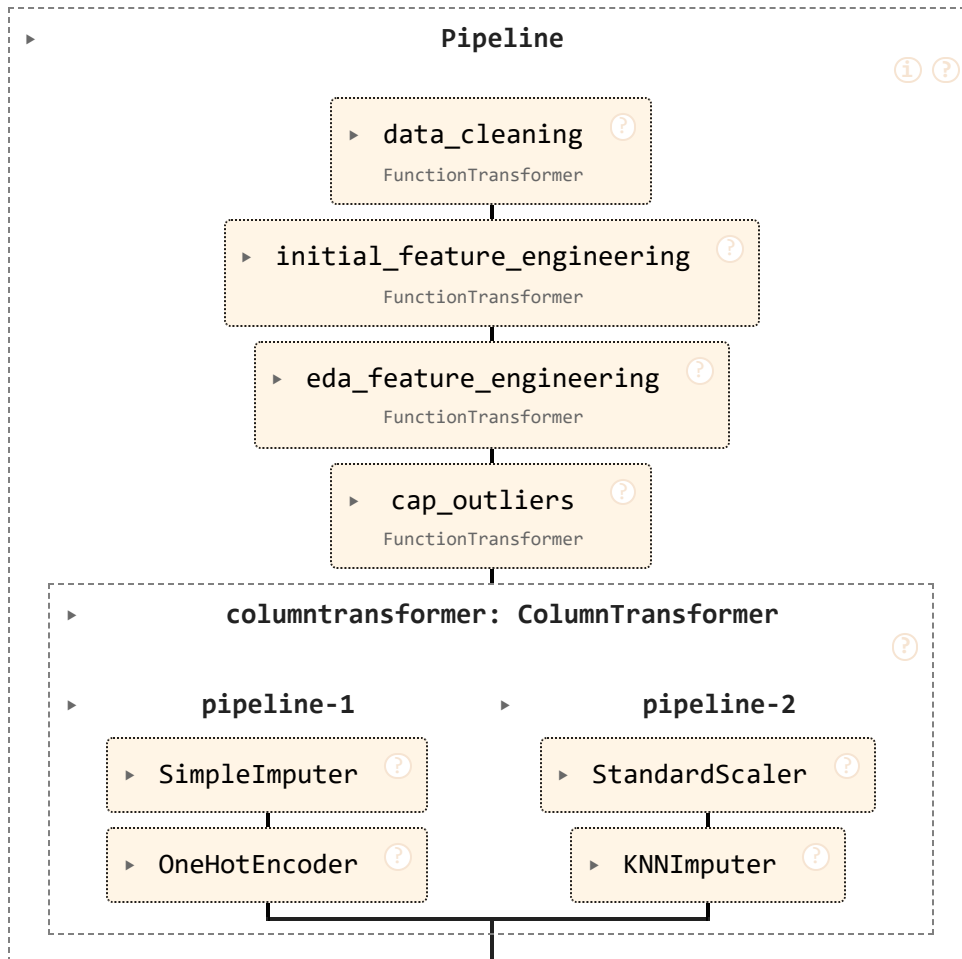
### 5.5 Create an entire pipeline for all preprocessing steps!¶

Creating a comprehensive preprocessing pipeline for ML is essential for consistency, efficiency, and reproducibility. It prevents data leakage, simplifies scaling, and integrates hyperparameter tuning seamlessly. Such a pipeline also aids in model deployment, enhancing performance, and maintaining a reliable ML workflow.

```
In [59]: pipe = make_pipeline(
            initial_clean,
            initial_feature_engineering,
            eda_feature_engineering,
            cap_outliers,
            impute_scale
        )
```

```
# Let's see how it looks
pipe
```

Out[59]:



```
In [62]: X_train_pp = pipe.fit_transform(X_train)
```

## 6. Modeling

We'll start by exploring models for potential strong performance. First, we'll evaluate them using cross-validation with stratified folds to maintain class proportions. The goal is to identify promising models before fine-tuning hyperparameters.

### Let's Remember Our Initial Target:

"The CEO, in particular, has given a ballpark of the target lead conversion rate to be around 80."

So, with this in mind, we can select our most important performance measure. In this case, we want to ensure that a high percentage of predicted leads convert to a customer, which means we're looking for a high precision score.

Does that mean we won't care about potential leads not detected (Low recall)?

Not at all. If we tune our models to optimize only for precision, we might be very accurate in their positive predictions but miss many actual positive cases. This translates into leaving money on the table—potential customers that won't convert.

## Display function and StratifiedKFold

```
In [63]: # Use stratified fold for ensure that we shuffle the dataset and conserve classes
skfold = StratifiedKFold(5, shuffle=True, random_state=12)

def display_scores(model,scores,pred):
    print(f'----- {model} -----')
    print('')
    print("----- Cross validation scores:")
    print("Scores:", scores)
    print("Mean:", scores.mean())
    print("Standard deviation:", scores.std())
    print('')
    print("----- Scores in the training set:")
    print("Precision:", precision_score(y_train,pred))
    print("Recall:", recall_score(y_train,pred))
    print("F1 score:", f1_score(y_train,pred))
    print("ROC - AUC score:", roc_auc_score(y_train,pred))
```

## 6.1 Logistic Regression

```
In [64]: lr = LogisticRegression()
lr_scores = cross_val_score(lr, X_train_pp, y_train,
                           cv=skfold, scoring='f1')
lr.fit(X_train_pp,y_train)
lr_pred = lr.predict(X_train_pp)

# Precision and recall curve
lr_prec, lr_recall, lr_threshold = precision_recall_curve(y_train, lr_pred, pos_lab
lr_prdisplay = PrecisionRecallDisplay(precision=lr_prec, recall=lr_recall)

# Display Scores
display_scores('Logistic Regression',lr_scores,lr_pred)
```

----- Logistic Regression -----

----- Cross validation scores:

Scores: [0.91741472 0.91756272 0.91785714 0.91896705 0.93109541]

Mean: 0.9205794094986535

Standard deviation: 0.0052860620627966275

----- Scores in the training set:

Precision: 0.9382491827097712

Recall: 0.9066339066339066

F1 score: 0.9221706533380936

ROC - AUC score: 0.9346068498610871

## 6.2 Support Vector Machine

```
In [65]: svc = SVC()
svc_scores = cross_val_score(svc, X_train_pp, y_train,
                             cv=skfold, scoring='f1')
svc.fit(X_train_pp, y_train)
svc_pred = svc.predict(X_train_pp)

# Precision and recall curve
svc_prec, svc_recall, svc_threshold = precision_recall_curve(y_train, svc_pred, pos
svc_prdisplay = PrecisionRecallDisplay(precision=svc_prec, recall=svc_recall)

# Display scores
display_scores('Support Vector Machine',svc_scores,svc_pred)
```

----- Support Vector Machine -----

----- Cross validation scores:

Scores: [0.91838565 0.92184725 0.9341637 0.91954023 0.93848858]

Mean: 0.9264850809036906

Standard deviation: 0.008226619600733422

----- Scores in the training set:

Precision: 0.9437074220150592

Recall: 0.9238329238329238

F1 score: 0.9336644200070947

ROC - AUC score: 0.9446371310778091

```
In [66]: ### 6.3 Decission Trees
```

```
In [67]: tree = DecisionTreeClassifier(random_state = 7)
tree_scores = cross_val_score(tree, X_train_pp, y_train,
                              cv=skfold, scoring='f1')
tree.fit(X_train_pp, y_train)
tree_pred = tree.predict(X_train_pp)

# Precision and recall curve
tree_prec, tree_recall, tree_threshold = precision_recall_curve(y_train, tree_pred,
tree_prdisplay = PrecisionRecallDisplay(precision=tree_prec, recall=tree_recall)

# Display scores
display_scores('Decission Tree',tree_scores,tree_pred)
```

----- Decission Tree -----

----- Cross validation scores:

Scores: [0.89492119 0.89612676 0.89806678 0.89137931 0.89837746]

Mean: 0.8957743001598371

Standard deviation: 0.002537709857476365

----- Scores in the training set:

Precision: 0.9912434325744308

Recall: 0.9933309933309933

F1 score: 0.9922861150070126

ROC - AUC score: 0.9939140108631633

```
In [68]: ### 6.4 Random Forest
```



```
In [69]: rf = RandomForestClassifier(random_state=10,
                                   oob_score=True)
rf_scores = cross_val_score(rf, X_train_pp, y_train,
                           cv=skfold, scoring='f1')
rf.fit(X_train_pp, y_train)
rf_pred = rf.predict(X_train_pp)
rf_pred_proba = rf.predict_proba(X_train_pp)

# Precision and recall curve
rf_prec, rf_recall, rf_threshold = precision_recall_curve(y_train, rf_pred_proba[:,
rf_prdisplay = PrecisionRecallDisplay(precision=rf_prec, recall=rf_recall)

# Display scores
display_scores('Random Forest', rf_scores, rf_pred)
print('Oob score: ', rf.oob_score_)
```

----- Random Forest -----

----- Cross validation scores:

Scores: [0.91921005 0.91974752 0.93027361 0.9215859 0.93960924]

Mean: 0.9260852646703863

Standard deviation: 0.007850093308105268

----- Scores in the training set:

Precision: 0.9908995449772489

Recall: 0.9936819936819937

F1 score: 0.9922888187872415

ROC - AUC score: 0.9939794516065702

Oob score: 0.9422348484848485

## 6.5 Gradient Boosting

```
In [70]: xg = GradientBoostingClassifier(random_state=11)
xg_scores = cross_val_score(xg, X_train_pp, y_train,
                           cv=skfold, scoring='f1')
xg.fit(X_train_pp, y_train)
xg_pred = xg.predict(X_train_pp)

# Precision and recall curve
xg_prec, xg_recall, xg_threshold = precision_recall_curve(y_train, xg_pred, pos_lab
xg_prdisplay = PrecisionRecallDisplay(precision=xg_prec, recall=xg_recall)

# Display scores
display_scores('Gradient Boosting', xg_scores, xg_pred)
```

----- Gradient Boosting -----

----- Cross validation scores:

Scores: [0.91862568 0.92196007 0.9309417 0.92197309 0.94201606]

Mean: 0.927103321202558

Standard deviation: 0.008506063041899237

----- Scores in the training set:

Precision: 0.9571322985957132

Recall: 0.9090909090909091

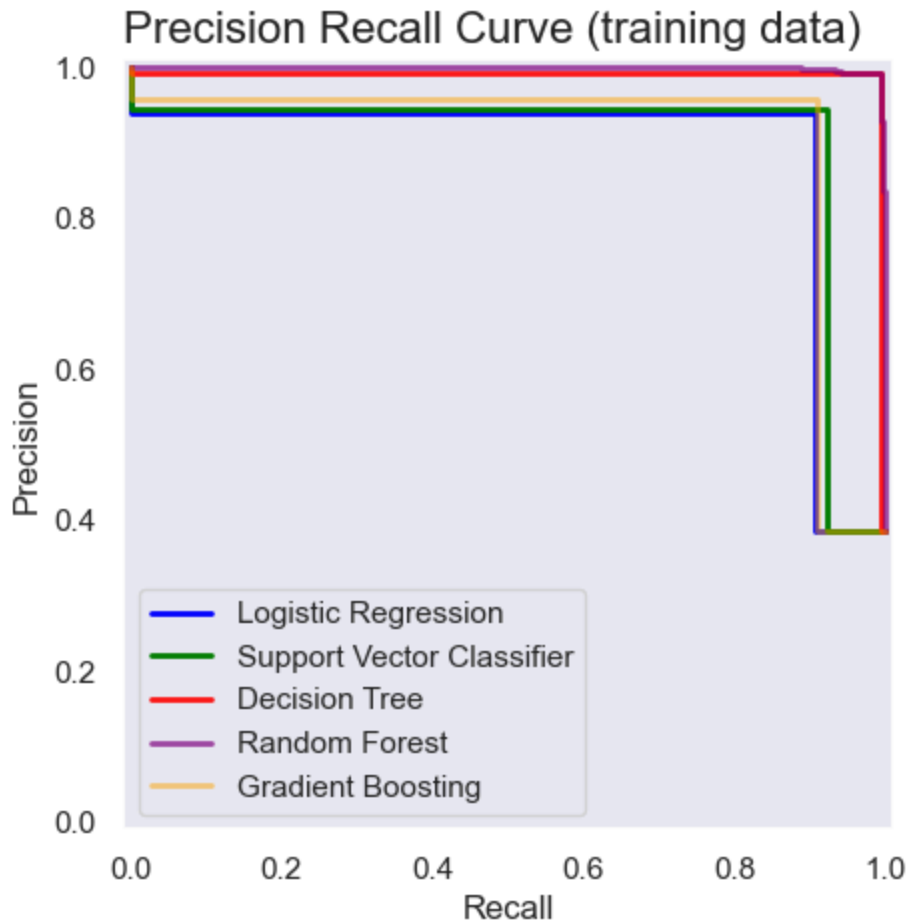
F1 score: 0.9324932493249325

ROC - AUC score: 0.9417785604226282

## 7. Select the best model and tune them

### 7.1 Recall - Precision Curve for each model

```
In [71]: import matplotlib.pyplot as plt
fig, ax = plt.subplots(figsize=(8,5))
lr_prdisplay.plot(ax=ax, label='Logistic Regression', color='blue', linewidth=2)
svc_prdisplay.plot(ax=ax, label='Support Vector Classifier', color='green', linewidth=2)
tree_prdisplay.plot(ax=ax, label='Decision Tree', color='red', linewidth=2, alpha=.5)
rf_prdisplay.plot(ax=ax, label='Random Forest', color='purple', linewidth=2, alpha=.5)
xg_prdisplay.plot(ax=ax, label='Gradient Boosting', color='orange', linewidth=2, alpha=.5)
plt.title('Precision Recall Curve (training data)', size=16, loc='left')
plt.show()
```



## 8. Make Our Prediction

At this point we are already:

1.

Completed the entire data preprocessing and exploration.#### 2.

We exclusively used the training dataset to eliminate any potential human bias#### 3. .

Additionally, we've incorporated all the preprocessing steps into a pipeline to prevent any data leak.

4. e.

Next, we selected the most promising models (without tuning) and applied cross-validation to assess their performan#### 5. ce. Following that, we fine-tuned those models using RandomizedSearchCV and identified the best on

### 8.2 Apply al the preprocessing pipeline to the test dataset

```
In [72]: X_test = test.drop(columns=['Converted']) # More readable way to drop a column
y_test = test['Converted'] # Directly selecting the column
```

```
# Viewing the first row as a NumPy array
X_test.iloc[:1].to_numpy()
```

```
Out[72]: array([[ 'b4d86fa1-53d9-4a27-8d0c-f6603a562184', 634844, 'API', 'Google',
                'No', 'No', 2.0, 1551, 1.0, 'SMS Sent', 'India', nan, nan, nan,
                nan, 'No', 'No', 'No', 'No', 'No', 'No', 'No', 'No', nan, nan,
                'No', 'No', nan, nan, '02.Medium', '02.Medium', 15.0, 15.0, 'No',
                'No', 'SMS Sent']], dtype=object)
```

```
In [73]: # apply all the preprocessing steps to the test dataset
X_test_pp = pipe.transform(X_test)
X_test_pp.toarray()[:1]
```

```
Out[73]: array([[ 1.          ,  0.          ,  0.          ,  0.          ,  0.          ,
                  0.          ,  1.          ,  0.          ,  0.          ,  0.          ,
                  0.          ,  0.          ,  0.          ,  1.          ,  0.          ,
                  0.          ,  0.          ,  0.          ,  0.          ,  0.          ,
                  0.          ,  1.          ,  0.          ,  0.          ,  0.          ,
                  0.          ,  0.          ,  0.          ,  1.          ,  0.          ,
                  0.          ,  0.          ,  0.          ,  0.          ,  0.          ,
                  0.          ,  0.          ,  1.          ,  0.          ,  0.          ,
                  0.          ,  0.          ,  0.          ,  0.          ,  0.          ,
                  0.          ,  1.          ,  0.          ,  1.          ,  0.          ,
                  0.          ,  0.          ,  0.          ,  0.          ,  0.          ,
                  0.          ,  0.          ,  0.          ,  0.          ,  0.          ,
                  0.          ,  0.          ,  0.          ,  0.          ,  0.          ,
                  0.          ,  0.          ,  1.          ,  0.          ,  0.          ,
                  0.          , -0.33742588, -0.62669473,  1.93925017,  0.50806717,
                  -0.74681403,  1.33195148]])
```

## 10.2 Random Forest with hyperparameter tuned

```
In [74]: print(X_train.dtypes)
from sklearn.preprocessing import OneHotEncoder
import pandas as pd

# Apply One-Hot Encoding
X_train = pd.get_dummies(X_train, drop_first=True)
X_test = pd.get_dummies(X_test, drop_first=True)

# Ensure both datasets have the same columns
X_train, X_test = X_train.align(X_test, join="left", axis=1, fill_value=0)

from sklearn.preprocessing import LabelEncoder

encoder = LabelEncoder()
for col in X_train.select_dtypes(include=['object']).columns:
    X_train[col] = encoder.fit_transform(X_train[col])
    X_test[col] = encoder.transform(X_test[col])

X_train = X_train.fillna(0) # Replace NaN with 0
X_test = X_test.fillna(0)

rf_randomcv.fit(X_train, y_train)
```

```

Prospect ID          object
Lead Number          int64
Lead Origin           object
Lead Source           object
Do Not Email         object
Do Not Call           object
TotalVisits           float64
Total Time Spent on Website int64
Page Views Per Visit  float64
Last Activity         object
Country              object
Specialization        object
How did you hear about X Education object
What is your current occupation object
What matters most to you in choosing a course object
Search               object
Magazine              object
Newspaper Article     object
X Education Forums    object
Newspaper             object
Digital Advertisement object
Through Recommendations object
Receive More Updates About Our Courses object
Tags                 object
Lead Quality          object
Update me on Supply Chain Content object
Get updates on DM Content object
Lead Profile          object
City                 object
Asymmetrique Activity Index object
Asymmetrique Profile Index object
Asymmetrique Activity Score float64
Asymmetrique Profile Score float64
I agree to pay the amount through cheque object
A free copy of Mastering The Interview object
Last Notable Activity object
dtype: object

```

```

-----
NameError                                Traceback (most recent call last)
Cell In[74], line 22
     19 X_train = X_train.fillna(0) # Replace NaN with 0
     20 X_test = X_test.fillna(0)
--> 22 rf_randomcv.fit(X_train, y_train)

NameError: name 'rf_randomcv' is not defined

```

```

In [ ]: rf_rcv_pred = rf_randomcv.predict(X_test_pp)
print("Precision:", precision_score(y_test, rf_rcv_pred))
print("Recall:", recall_score(y_test, rf_rcv_pred))
print("F1 score:", f1_score(y_test, rf_rcv_pred))
print("ROC - AUC score:", roc_auc_score(y_test, rf_rcv_pred))

```

## 10.3 Random Forest without hyperparameter tuned

```
In [ ]: rf_pred_test = rf.predict(X_test_pp)
print("Precision:", precision_score(y_test, rf_pred_test))
print("Recall:", recall_score(y_test, rf_pred_test))
print("F1 score:", f1_score(y_test, rf_pred_test))
print("ROC - AUC score:", roc_auc_score(y_test, rf_pred_test))
```

## 10.4 Let's plot the confusion matrix for both models!

```
In [ ]: fig, ax = plt.subplots(1, 2, figsize=(10, 4))

# Random Forest tuned
cm1 = confusion_matrix(y_test, rf_rcv_pred)
sns.heatmap(cm1, annot=True, fmt='d', cmap='Greens', ax = ax[0], cbar=False)
ax[0].xaxis.set_ticklabels(['Not converted', 'Converted'])
ax[0].yaxis.set_ticklabels(['Not converted', 'Converted'])
ax[0].set_title('RF with hyperparameters tuning', loc='left')
ax[0].set_xlabel('Predicted')
ax[0].set_ylabel('True')

# Random Forest without tuning
cm2 = confusion_matrix(y_test, rf_pred_test)
sns.heatmap(cm2, annot=True, fmt='d', cmap='Blues', ax=ax[1], cbar=False)
ax[1].xaxis.set_ticklabels(['Not converted', 'Converted'])
ax[1].yaxis.set_ticklabels(['Not converted', 'Converted'])
ax[1].set_title('RF without hyperparameters tuning', loc='left')
ax[1].set_xlabel('Predicted')
ax[1].set_ylabel('True')

plt.tight_layout()
plt.show()
```

```
In [ ]: import matplotlib.pyplot as plt

# Plot a histogram of Lead scores
plt.hist(lead_scoring, bins=20, edgecolor='black')
plt.xlabel("Lead Score")
plt.ylabel("Frequency")
plt.title("Distribution of Lead Scores")
plt.show()
```

```
In [ ]: plt.scatter(lead_prediction, lead_scoring, alpha=0.5)
plt.xlabel("Lead Prediction (0 or 1)")
plt.ylabel("Lead Score")
plt.title("Lead Prediction vs. Lead Score")
plt.show()
```

## Submission

- Class predictions in the left, and probabilities to convert into a customer on the right.

```
In [ ]: import numpy as np
```

```
# Assuming rf_randomcv is a trained RandomForest model and X_test_pp is preprocessed
lead_scoring = rf_randomcv.predict_proba(X_test_pp)[: , 1] # Getting probabilities
lead_prediction = rf_randomcv.predict(X_test_pp) # Getting predictions

# Combining predictions and probabilities
results = np.round(np.c_[lead_prediction, lead_scoring], 2)

# Display the first 10 rows
print(results[:10])
```

## Conclusion

In summary, our data science project focused on fine-tuning lead scoring for X Education. We aimed to exceed an 80% precision goal, which we not only met but exceeded. Throughout our journey, we identified key factors like phone interactions, referrals, and online engagement that strongly correlated with lead conversion, leading to actionable strategies.

One notable achievement was the development of an automated lead scoring algorithm that not only improved lead assessment precision but also streamlined operational efficiency. By targeting promising leads, X Education could reduce sales team costs significantly.

Our journey involved thorough data exploration, preprocessing, and model development, ensuring consistency and mitigating bias. We systematically evaluated models, with the tuned Random Forest model achieving an impressive F1 score of 0.9287 and a precision score of 0.9527 on the test dataset.

This data-driven journey provides X Education with actionable insights to enhance efficiency and revenue growth, positioning the company for a transformative phase.

If you've read until here, thank you. I hope you found this information helpful and interesting in some way. Your feedback is greatly appreciated. Best regards.